



Language Specification

Version 3.0

Copyright © Microsoft Corporation 1999-2007. All Rights Reserved.
Please send corrections, comments, and other feedback to sharp@microsoft.com

Notice

© 1999-2007 Microsoft Corporation. All rights reserved.

Microsoft, Windows, Visual Basic, Visual C#, and Visual C++ are either registered trademarks or trademarks of Microsoft Corporation in the U.S.A. and/or other countries/regions.

Other product and company names mentioned herein may be the trademarks of their respective owners.

Copyright © Microsoft Corporation 1999-2007. All Rights Reserved.

Table of Contents

Table of Contents

1. Introduction.....	1
1.1 Hello world.....	1
1.2 Program structure.....	2
1.3 Types and variables.....	4
1.4 Expressions.....	6
1.5 Statements.....	8
1.6 Classes and objects.....	12
1.6.1 Members.....	12
1.6.2 Accessibility.....	13
1.6.3 Type parameters.....	13
1.6.4 Base classes.....	14
1.6.5 Fields.....	14
1.6.6 Methods.....	15
1.6.6.1 Parameters.....	15
1.6.6.2 Method body and local variables.....	16
1.6.6.3 Static and instance methods.....	17
1.6.6.4 Virtual, override, and abstract methods.....	18
1.6.6.5 Method overloading.....	20
1.6.7 Other function members.....	21
1.6.7.1 Constructors.....	22
1.6.7.2 Properties.....	23
1.6.7.3 Indexers.....	23
1.6.7.4 Events.....	24
1.6.7.5 Operators.....	24
1.6.7.6 Destructors.....	25
1.7 Structs.....	25
1.8 Arrays.....	26
1.9 Interfaces.....	27
1.10 Enums.....	29
1.11 Delegates.....	30
1.12 Attributes.....	31
2. Lexical structure.....	33
2.1 Programs.....	33
2.2 Grammars.....	33
2.2.1 Grammar notation.....	33
2.2.2 Lexical grammar.....	34
2.2.3 Syntactic grammar.....	34
2.3 Lexical analysis.....	34
2.3.1 Line terminators.....	35
2.3.2 Comments.....	35
2.3.3 White space.....	37
2.4 Tokens.....	37
2.4.1 Unicode character escape sequences.....	37
2.4.2 Identifiers.....	38
2.4.3 Keywords.....	40
2.4.4 Literals.....	40
2.4.4.1 Boolean literals.....	40
2.4.4.2 Integer literals.....	40

C# LANGUAGE SPECIFICATION

2.4.4.3 Real literals.....	42
2.4.4.4 Character literals.....	42
2.4.4.5 String literals.....	43
2.4.4.6 The null literal.....	45
2.4.5 Operators and punctuators.....	45
2.5 Pre-processing directives.....	45
2.5.1 Conditional compilation symbols.....	47
2.5.2 Pre-processing expressions.....	47
2.5.3 Declaration directives.....	48
2.5.4 Conditional compilation directives.....	49
2.5.5 Diagnostic directives.....	51
2.5.6 Region directives.....	52
2.5.7 Line directives.....	52
2.5.8 Pragma directives.....	53
2.5.8.1 Pragma warning.....	53
3. Basic concepts.....	55
3.1 Application Startup.....	55
3.2 Application termination.....	56
3.3 Declarations.....	56
3.4 Members.....	58
3.4.1 Namespace members.....	58
3.4.2 Struct members.....	59
3.4.3 Enumeration members.....	59
3.4.4 Class members.....	59
3.4.5 Interface members.....	60
3.4.6 Array members.....	60
3.4.7 Delegate members.....	60
3.5 Member access.....	60
3.5.1 Declared accessibility.....	60
3.5.2 Accessibility domains.....	61
3.5.3 Protected access for instance members.....	63
3.5.4 Accessibility constraints.....	64
3.6 Signatures and overloading.....	65
3.7 Scopes.....	66
3.7.1 Name hiding.....	69
3.7.1.1 Hiding through nesting.....	69
3.7.1.2 Hiding through inheritance.....	70
3.8 Namespace and type names.....	71
3.8.1 Fully qualified names.....	73
3.9 Automatic memory management.....	73
3.10 Execution order.....	76
4. Types.....	77
4.1 Value types.....	77
4.1.1 The System.ValueType type.....	78
4.1.2 Default constructors.....	78
4.1.3 Struct types.....	79
4.1.4 Simple types.....	79
4.1.5 Integral types.....	80
4.1.6 Floating point types.....	81

Table of Contents

4.1.7 The decimal type.....	82
4.1.8 The bool type.....	83
4.1.9 Enumeration types.....	83
4.1.10 Nullable types.....	83
4.2 Reference types.....	84
4.2.1 Class types.....	84
4.2.2 The object type.....	85
4.2.3 The string type.....	85
4.2.4 Interface types.....	85
4.2.5 Array types.....	85
4.2.6 Delegate types.....	85
4.3 Boxing and unboxing.....	86
4.3.1 Boxing conversions.....	86
4.3.2 Unboxing conversions.....	87
4.4 Constructed types.....	88
4.4.1 Type arguments.....	89
4.4.2 Open and closed types.....	89
4.4.3 Bound and unbound types.....	89
4.4.4 Satisfying constraints.....	89
4.5 Type parameters.....	90
4.6 Expression tree types.....	91
5. Variables.....	93
5.1 Variable categories.....	93
5.1.1 Static variables.....	93
5.1.2 Instance variables.....	93
5.1.2.1 Instance variables in classes.....	94
5.1.2.2 Instance variables in structs.....	94
5.1.3 Array elements.....	94
5.1.4 Value parameters.....	94
5.1.5 Reference parameters.....	94
5.1.6 Output parameters.....	95
5.1.7 Local variables.....	95
5.2 Default values.....	96
5.3 Definite assignment.....	96
5.3.1 Initially assigned variables.....	97
5.3.2 Initially unassigned variables.....	97
5.3.3 Precise rules for determining definite assignment.....	97
5.3.3.1 General rules for statements.....	98
5.3.3.2 Block statements, checked, and unchecked statements.....	98
5.3.3.3 Expression statements.....	98
5.3.3.4 Declaration statements.....	99
5.3.3.5 If statements.....	99
5.3.3.6 Switch statements.....	99
5.3.3.7 While statements.....	99
5.3.3.8 Do statements.....	100
5.3.3.9 For statements.....	100
5.3.3.10 Break, continue, and goto statements.....	100
5.3.3.11 Throw statements.....	100
5.3.3.12 Return statements.....	100
5.3.3.13 Try-catch statements.....	101

C# LANGUAGE SPECIFICATION

5.3.3.14 Try-finally statements.....	101
5.3.3.15 Try-catch-finally statements.....	101
5.3.3.16 Foreach statements.....	102
5.3.3.17 Using statements.....	102
5.3.3.18 Lock statements.....	103
5.3.3.19 Yield statements.....	103
5.3.3.20 General rules for simple expressions.....	103
5.3.3.21 General rules for expressions with embedded expressions.....	103
5.3.3.22 Invocation expressions and object creation expressions.....	104
5.3.3.23 Simple assignment expressions.....	104
5.3.3.24 && expressions.....	104
5.3.3.25 expressions.....	105
5.3.3.26 ! expressions.....	106
5.3.3.27 ?? expressions.....	106
5.3.3.28 ? expressions.....	107
5.3.3.29 Anonymous functions.....	107
5.4 Variable references.....	108
5.5 Atomicity of variable references.....	108
6. Conversions.....	109
6.1 Implicit conversions.....	109
6.1.1 Identity conversion.....	109
6.1.2 Implicit numeric conversions.....	109
6.1.3 Implicit enumeration conversions.....	110
6.1.4 Implicit nullable conversions.....	110
6.1.5 Null literal conversions.....	110
6.1.6 Implicit reference conversions.....	111
6.1.7 Boxing conversions.....	111
6.1.8 Implicit constant expression conversions.....	112
6.1.9 Implicit conversions involving type parameters.....	112
6.1.10 User-defined implicit conversions.....	112
6.1.11 Anonymous function conversions and method group conversions.....	112
6.2 Explicit conversions.....	112
6.2.1 Explicit numeric conversions.....	113
6.2.2 Explicit enumeration conversions.....	114
6.2.3 Explicit nullable conversions.....	115
6.2.4 Explicit reference conversions.....	115
6.2.5 Unboxing conversions.....	116
6.2.6 Explicit conversions involving type parameters.....	116
6.2.7 User-defined explicit conversions.....	117
6.3 Standard conversions.....	117
6.3.1 Standard implicit conversions.....	117
6.3.2 Standard explicit conversions.....	118
6.4 User-defined conversions.....	118
6.4.1 Permitted user-defined conversions.....	118
6.4.2 Lifted conversion operators.....	118
6.4.3 Evaluation of user-defined conversions.....	118
6.4.4 User-defined implicit conversions.....	119
6.4.5 User-defined explicit conversions.....	120
6.5 Anonymous function conversions.....	121
6.5.1 Evaluation of anonymous function conversions to delegate types.....	122

Table of Contents

6.5.2 Evaluation of anonymous function conversions to expression tree types.....	123
6.5.3 Implementation example.....	123
6.6 Method group conversions.....	125
7. Expressions.....	129
7.1 Expression classifications.....	129
7.1.1 Values of expressions.....	130
7.2 Operators.....	130
7.2.1 Operator precedence and associativity.....	130
7.2.2 Operator overloading.....	131
7.2.3 Unary operator overload resolution.....	133
7.2.4 Binary operator overload resolution.....	133
7.2.5 Candidate user-defined operators.....	133
7.2.6 Numeric promotions.....	134
7.2.6.1 Unary numeric promotions.....	134
7.2.6.2 Binary numeric promotions.....	134
7.2.7 Lifted operators.....	135
7.3 Member lookup.....	136
7.3.1 Base types.....	137
7.4 Function members.....	137
7.4.1 Argument lists.....	140
7.4.2 Type inference.....	142
7.4.2.1 The first phase.....	143
7.4.2.2 The second phase.....	143
7.4.2.3 Input types.....	143
7.4.2.4 Output types.....	143
7.4.2.5 Dependence.....	143
7.4.2.6 Output type inferences.....	143
7.4.2.7 Explicit parameter type inferences.....	144
7.4.2.8 Exact inferences.....	144
7.4.2.9 Lower-bound inferences.....	144
7.4.2.10 Fixing.....	144
7.4.2.11 Inferred return type.....	145
7.4.2.12 Type inference for conversion of method groups.....	146
7.4.2.13 Finding the best common type of a set of expressions.....	146
7.4.3 Overload resolution.....	146
7.4.3.1 Applicable function member.....	147
7.4.3.2 Better function member.....	147
7.4.3.3 Better conversion from expression.....	148
7.4.3.4 Better conversion from type.....	149
7.4.3.5 Overloading in generic classes.....	149
7.4.4 Function member invocation.....	150
7.4.4.1 Invocations on boxed instances.....	151
7.5 Primary expressions.....	151
7.5.1 Literals.....	152
7.5.2 Simple names.....	152
7.5.2.1 Invariant meaning in blocks.....	154
7.5.3 Parenthesized expressions.....	154
7.5.4 Member access.....	155
7.5.4.1 Identical simple names and type names.....	156
7.5.4.2 Grammar ambiguities.....	157

C# LANGUAGE SPECIFICATION

7.5.5 Invocation expressions.....	158
7.5.5.1 Method invocations.....	158
7.5.5.2 Extension method invocations.....	159
7.5.5.3 Delegate invocations.....	161
7.5.6 Element access.....	162
7.5.6.1 Array access.....	162
7.5.6.2 Indexer access.....	163
7.5.7 This access.....	163
7.5.8 Base access.....	164
7.5.9 Postfix increment and decrement operators.....	164
7.5.10 The new operator.....	165
7.5.10.1 Object creation expressions.....	166
7.5.10.2 Object initializers.....	167
7.5.10.3 Collection initializers.....	169
7.5.10.4 Array creation expressions.....	170
7.5.10.5 Delegate creation expressions.....	172
7.5.10.6 Anonymous object creation expressions.....	173
7.5.11 The typeof operator.....	175
7.5.12 The checked and unchecked operators.....	176
7.5.13 Default value expressions.....	178
7.5.14 Anonymous method expressions.....	179
7.6 Unary operators.....	179
7.6.1 Unary plus operator.....	179
7.6.2 Unary minus operator.....	179
7.6.3 Logical negation operator.....	180
7.6.4 Bitwise complement operator.....	180
7.6.5 Prefix increment and decrement operators.....	181
7.6.6 Cast expressions.....	181
7.7 Arithmetic operators.....	182
7.7.1 Multiplication operator.....	182
7.7.2 Division operator.....	183
7.7.3 Remainder operator.....	184
7.7.4 Addition operator.....	185
7.7.5 Subtraction operator.....	187
7.8 Shift operators.....	189
7.9 Relational and type-testing operators.....	190
7.9.1 Integer comparison operators.....	190
7.9.2 Floating-point comparison operators.....	191
7.9.3 Decimal comparison operators.....	192
7.9.4 Boolean equality operators.....	192
7.9.5 Enumeration comparison operators.....	192
7.9.6 Reference type equality operators.....	192
7.9.7 String equality operators.....	194
7.9.8 Delegate equality operators.....	194
7.9.9 Equality operators and null.....	195
7.9.10 The is operator.....	195
7.9.11 The as operator.....	196
7.10 Logical operators.....	196
7.10.1 Integer logical operators.....	197
7.10.2 Enumeration logical operators.....	197
7.10.3 Boolean logical operators.....	197

Table of Contents

7.10.4 Nullable boolean logical operators	197
7.11 Conditional logical operators.....	198
7.11.1 Boolean conditional logical operators.....	199
7.11.2 User-defined conditional logical operators	199
7.12 The null coalescing operator.....	199
7.13 Conditional operator.....	200
7.14 Anonymous function expressions.....	201
7.14.1 Anonymous function signatures	203
7.14.2 Anonymous function bodies	203
7.14.3 Overload resolution.....	204
7.14.4 Outer variables	204
7.14.4.1 Captured outer variables.....	204
7.14.4.2 Instantiation of local variables	205
7.14.5 Evaluation of anonymous function expressions.....	207
7.15 Query expressions	207
7.15.1 Ambiguities in query expressions.....	209
7.15.2 Query expression translation.....	209
7.15.2.1 Select and groupby clauses with continuations.....	209
7.15.2.2 Explicit range variable types	210
7.15.2.3 Degenerate query expressions.....	210
7.15.2.4 From, let, where, join and orderby clauses.....	211
7.15.2.5 Select clauses.....	214
7.15.2.6 Groupby clauses	214
7.15.2.7 Transparent identifiers.....	215
7.15.3 The query expression pattern.....	216
7.16 Assignment operators	217
7.16.1 Simple assignment	218
7.16.2 Compound assignment.....	220
7.16.3 Event assignment	221
7.17 Expression	221
7.18 Constant expressions	221
7.19 Boolean expressions	223
8. Statements.....	225
8.1 End points and reachability	225
8.2 Blocks	227
8.2.1 Statement lists.....	227
8.3 The empty statement.....	228
8.4 Labeled statements	228
8.5 Declaration statements.....	229
8.5.1 Local variable declarations	229
8.5.2 Local constant declarations.....	230
8.6 Expression statements	231
8.7 Selection statements	231
8.7.1 The if statement	231
8.7.2 The switch statement.....	232
8.8 Iteration statements.....	235
8.8.1 The while statement	236
8.8.2 The do statement.....	236
8.8.3 The for statement	237
8.8.4 The foreach statement	238

C# LANGUAGE SPECIFICATION

8.9 Jump statements.....	240
8.9.1 The break statement	241
8.9.2 The continue statement	242
8.9.3 The goto statement.....	242
8.9.4 The return statement	243
8.9.5 The throw statement	244
8.10 The try statement	245
8.11 The checked and unchecked statements	248
8.12 The lock statement	248
8.13 The using statement.....	249
8.14 The yield statement.....	251
9. Namespaces.....	253
9.1 Compilation units	253
9.2 Namespace declarations	253
9.3 Extern aliases	255
9.4 Using directives	255
9.4.1 Using alias directives.....	256
9.4.2 Using namespace directives.....	258
9.5 Namespace members.....	260
9.6 Type declarations.....	260
9.7 Namespace alias qualifiers	261
9.7.1 Uniqueness of aliases.....	262
10. Classes	263
10.1 Class declarations	263
10.1.1 Class modifiers	263
10.1.1.1 Abstract classes	264
10.1.1.2 Sealed classes	264
10.1.1.3 Static classes	264
10.1.2 Partial modifier	265
10.1.3 Type parameters	265
10.1.4 Class base specification	266
10.1.4.1 Base classes	266
10.1.4.2 Interface implementations	267
10.1.5 Type parameter constraints	268
10.1.6 Class body.....	272
10.2 Partial types	272
10.2.1 Attributes	272
10.2.2 Modifiers.....	272
10.2.3 Type parameters and constraints.....	273
10.2.4 Base class	273
10.2.5 Base interfaces	273
10.2.6 Members	274
10.2.7 Partial methods	274
10.2.8 Name binding.....	277
10.3 Class members.....	277
10.3.1 The instance type	279
10.3.2 Members of constructed types	279
10.3.3 Inheritance	280
10.3.4 The new modifier.....	281

Table of Contents

10.3.5 Access modifiers	281
10.3.6 Constituent types	281
10.3.7 Static and instance members	281
10.3.8 Nested types	282
10.3.8.1 Fully qualified name	283
10.3.8.2 Declared accessibility	283
10.3.8.3 Hiding	283
10.3.8.4 this access	284
10.3.8.5 Access to private and protected members of the containing type	285
10.3.8.6 Nested types in generic classes	286
10.3.9 Reserved member names	286
10.3.9.1 Member names reserved for properties	287
10.3.9.2 Member names reserved for events	287
10.3.9.3 Member names reserved for indexers	287
10.3.9.4 Member names reserved for destructors	288
10.4 Constants	288
10.5 Fields	289
10.5.1 Static and instance fields	291
10.5.2 Readonly fields	291
10.5.2.1 Using static readonly fields for constants	292
10.5.2.2 Versioning of constants and static readonly fields	292
10.5.3 Volatile fields	293
10.5.4 Field initialization	294
10.5.5 Variable initializers	294
10.5.5.1 Static field initialization	295
10.5.5.2 Instance field initialization	296
10.6 Methods	297
10.6.1 Method parameters	298
10.6.1.1 Value parameters	299
10.6.1.2 Reference parameters	300
10.6.1.3 Output parameters	301
10.6.1.4 Parameter arrays	301
10.6.2 Static and instance methods	304
10.6.3 Virtual methods	304
10.6.4 Override methods	306
10.6.5 Sealed methods	308
10.6.6 Abstract methods	309
10.6.7 External methods	310
10.6.8 Partial methods	311
10.6.9 Extension methods	311
10.6.10 Method body	311
10.6.11 Method overloading	312
10.7 Properties	312
10.7.1 Static and instance properties	313
10.7.2 Accessors	314
10.7.3 Automatically implemented properties	318
10.7.4 Accessibility	319
10.7.5 Virtual, sealed, override, and abstract accessors	320
10.8 Events	322
10.8.1 Field-like events	324
10.8.2 Event accessors	325

C# LANGUAGE SPECIFICATION

10.8.3 Static and instance events	326
10.8.4 Virtual, sealed, override, and abstract accessors	327
10.9 Indexers	327
10.9.1 Indexer overloading	331
10.10 Operators	331
10.10.1 Unary operators	333
10.10.2 Binary operators	333
10.10.3 Conversion operators	334
10.11 Instance constructors	336
10.11.1 Constructor initializers	337
10.11.2 Instance variable initializers	338
10.11.3 Constructor execution	338
10.11.4 Default constructors	340
10.11.5 Private constructors	340
10.11.6 Optional instance constructor parameters	341
10.12 Static constructors	341
10.13 Destructors	343
10.14 Iterators	345
10.14.1 Enumerator interfaces	345
10.14.2 Enumerable interfaces	345
10.14.3 Yield type	345
10.14.4 Enumerator objects	345
10.14.4.1 The MoveNext method	346
10.14.4.2 The Current property	347
10.14.4.3 The Dispose method	347
10.14.5 Enumerable objects	347
10.14.5.1 The GetEnumerator method	348
10.14.6 Implementation example	348
11. Structs	355
11.1 Struct declarations	355
11.1.1 Struct modifiers	355
11.1.2 Partial modifier	356
11.1.3 Struct interfaces	356
11.1.4 Struct body	356
11.2 Struct members	356
11.3 Class and struct differences	356
11.3.1 Value semantics	357
11.3.2 Inheritance	358
11.3.3 Assignment	358
11.3.4 Default values	358
11.3.5 Boxing and unboxing	359
11.3.6 Meaning of this	360
11.3.7 Field initializers	361
11.3.8 Constructors	361
11.3.9 Destructors	362
11.3.10 Static constructors	362
11.4 Struct examples	362
11.4.1 Database integer type	362
11.4.2 Database boolean type	364
12. Arrays	367

Table of Contents

12.1 Array types	367
12.1.1 The System.Array type	368
12.1.2 Arrays and the generic IList interface	368
12.2 Array creation	368
12.3 Array element access	369
12.4 Array members	369
12.5 Array covariance	369
12.6 Array initializers	369
13. Interfaces	373
13.1 Interface declarations	373
13.1.1 Interface modifiers	373
13.1.2 Partial modifier	373
13.1.3 Base interfaces	374
13.1.4 Interface body	374
13.2 Interface members	374
13.2.1 Interface methods	376
13.2.2 Interface properties	376
13.2.3 Interface events	376
13.2.4 Interface indexers	376
13.2.5 Interface member access	377
13.3 Fully qualified interface member names	378
13.4 Interface implementations	379
13.4.1 Explicit interface member implementations	380
13.4.2 Uniqueness of implemented interfaces	382
13.4.3 Implementation of generic methods	383
13.4.4 Interface mapping	383
13.4.5 Interface implementation inheritance	386
13.4.6 Interface re-implementation	387
13.4.7 Abstract classes and interfaces	389
14. Enums	391
14.1 Enum declarations	391
14.2 Enum modifiers	391
14.3 Enum members	392
14.4 The System.Enum type	394
14.5 Enum values and operations	394
15. Delegates	395
15.1 Delegate declarations	395
15.2 Delegate compatibility	397
15.3 Delegate instantiation	397
15.4 Delegate invocation	398
16. Exceptions	401
16.1 Causes of exceptions	401
16.2 The System.Exception class	401
16.3 How exceptions are handled	401
16.4 Common Exception Classes	402
17. Attributes	405
17.1 Attribute classes	405

C# LANGUAGE SPECIFICATION

17.1.1 Attribute usage	405
17.1.2 Positional and named parameters	406
17.1.3 Attribute parameter types	407
17.2 Attribute specification	407
17.3 Attribute instances	412
17.3.1 Compilation of an attribute	413
17.3.2 Run-time retrieval of an attribute instance	413
17.4 Reserved attributes	413
17.4.1 The AttributeUsage attribute	413
17.4.2 The Conditional attribute	414
17.4.2.1 Conditional methods	414
17.4.2.2 Conditional attribute classes	416
17.4.3 The Obsolete attribute	417
17.5 Attributes for Interoperation	418
17.5.1 Interoperation with COM and Win32 components	419
17.5.2 Interoperation with other .NET languages	419
17.5.2.1 The IndexerName attribute	419
18. Unsafe code	421
18.1 Unsafe contexts	421
18.2 Pointer types	423
18.3 Fixed and moveable variables	426
18.4 Pointer conversions	426
18.5 Pointers in expressions	427
18.5.1 Pointer indirection	428
18.5.2 Pointer member access	428
18.5.3 Pointer element access	429
18.5.4 The address-of operator	430
18.5.5 Pointer increment and decrement	431
18.5.6 Pointer arithmetic	431
18.5.7 Pointer comparison	432
18.5.8 The sizeof operator	432
18.6 The fixed statement	433
18.7 Fixed size buffers	436
18.7.1 Fixed size buffer declarations	436
18.7.2 Fixed size buffers in expressions	438
18.7.3 Definite assignment checking	438
18.8 Stack allocation	439
18.9 Dynamic memory allocation	440
A. Documentation comments	443
A.1 Introduction	443
A.2 Recommended tags	444
A.2.1 <c>	445
A.2.2 <code>	445
A.2.3 <example>	446
A.2.4 <exception>	446
A.2.5 <include>	447
A.2.6 <list>	447
A.2.7 <para>	448
A.2.8 <param>	449

Table of Contents

A.2.9 <paramref>	449
A.2.10 <permission>	449
A.2.11 <summary>	450
A.2.12 <returns>	450
A.2.13 <see>	451
A.2.14 <seealso>	451
A.2.15 <summary>	452
A.2.16 <value>	452
A.2.17 <typeparam>	452
A.2.18 <typeparamref>	453
A.3 Processing the documentation file	453
A.3.1 ID string format	453
A.3.2 ID string examples	454
A.4 An example	458
A.4.1 C# source code	458
A.4.2 Resulting XML	460
B. Grammar	466
B.1 Lexical grammar	466
B.1.1 Line terminators	466
B.1.2 Comments	466
B.1.3 White space	467
B.1.4 Tokens	467
B.1.5 Unicode character escape sequences	467
B.1.6 Identifiers	468
B.1.7 Keywords	469
B.1.8 Literals	469
B.1.9 Operators and punctuators	471
B.1.10 Pre-processing directives	471
B.2 Syntactic grammar	474
B.2.1 Basic concepts	474
B.2.2 Types	474
B.2.3 Variables	475
B.2.4 Expressions	476
B.2.5 Statements	482
B.2.6 Namespaces	486
B.2.7 Classes	487
B.2.8 Structs	494
B.2.9 Arrays	494
B.2.10 Interfaces	495
B.2.11 Enums	496
B.2.12 Delegates	496
B.2.13 Attributes	497
B.3 Grammar extensions for unsafe code	498
C. References	503

1. Introduction

C# (pronounced “See Sharp”) is a simple, modern, object-oriented, and type-safe programming language. C# has its roots in the C family of languages and will be immediately familiar to C, C++, and Java programmers. C# is standardized by ECMA International as the *ECMA-334* standard and by ISO/IEC as the *ISO/IEC 23270* standard. Microsoft’s C# compiler for the .NET Framework is a conforming implementation of both of these standards.

C# is an object-oriented language, but C# further includes support for *component-oriented* programming. Contemporary software design increasingly relies on software components in the form of self-contained and self-describing packages of functionality. Key to such components is that they present a programming model with properties, methods, and events; they have attributes that provide declarative information about the component; and they incorporate their own documentation. C# provides language constructs to directly support these concepts, making C# a very natural language in which to create and use software components.

Several C# features aid in the construction of robust and durable applications: *Garbage collection* automatically reclaims memory occupied by unused objects; *exception handling* provides a structured and extensible approach to error detection and recovery; and the *type-safe* design of the language makes it impossible to read from uninitialized variables, to index arrays beyond their bounds, or to perform unchecked type casts.

C# has a *unified type system*. All C# types, including primitive types such as `int` and `double`, inherit from a single root `object` type. Thus, all types share a set of common operations, and values of any type can be stored, transported, and operated upon in a consistent manner. Furthermore, C# supports both user-defined reference types and value types, allowing dynamic allocation of objects as well as in-line storage of lightweight structures.

To ensure that C# programs and libraries can evolve over time in a compatible manner, much emphasis has been placed on *versioning* in C#’s design. Many programming languages pay little attention to this issue, and, as a result, programs written in those languages break more often than necessary when newer versions of dependent libraries are introduced. Aspects of C#’s design that were directly influenced by versioning considerations include the separate `virtual` and `override` modifiers, the rules for method overload resolution, and support for explicit interface member declarations.

The rest of this chapter describes the essential features of the C# language. Although later chapters describe rules and exceptions in a detail-oriented and sometimes mathematical manner, this chapter strives for clarity and brevity at the expense of completeness. The intent is to provide the reader with an introduction to the language that will facilitate the writing of early programs and the reading of later chapters.

1.1 Hello world

The “Hello, World” program is traditionally used to introduce a programming language. Here it is in C#:

```
using System;
class Hello
{
    static void Main() {
        Console.WriteLine("Hello, world");
    }
}
```

C# source files typically have the file extension `.cs`. Assuming that the “Hello, World” program is stored in the file `hello.cs`, the program can be compiled with the Microsoft C# compiler using the command line

C# LANGUAGE SPECIFICATION

```
csc hello.cs
```

which produces an executable assembly named `hello.exe`. The output produced by this application when it is run is

```
Hello, world
```

The “Hello, World” program starts with a `using` directive that references the `System` namespace. Namespaces provide a hierarchical means of organizing C# programs and libraries. Namespaces contain types and other namespaces—for example, the `System` namespace contains a number of types, such as the `Console` class referenced in the program, and a number of other namespaces, such as `IO` and `Collections`. A `using` directive that references a given namespace enables unqualified use of the types that are members of that namespace. Because of the `using` directive, the program can use `Console.WriteLine` as shorthand for `System.Console.WriteLine`.

The `Hello` class declared by the “Hello, World” program has a single member, the method named `Main`. The `Main` method is declared with the `static` modifier. While instance methods can reference a particular enclosing object instance using the keyword `this`, static methods operate without reference to a particular object. By convention, a static method named `Main` serves as the entry point of a program.

The output of the program is produced by the `WriteLine` method of the `Console` class in the `System` namespace. This class is provided by the .NET Framework class libraries, which, by default, are automatically referenced by the Microsoft C# compiler. Note that C# itself does not have a separate runtime library. Instead, the .NET Framework *is* the runtime library of C#.

1.2 Program structure

The key organizational concepts in C# are *programs*, *namespaces*, *types*, *members*, and *assemblies*. C# programs consist of one or more source files. Programs declare types, which contain members and can be organized into namespaces. Classes and interfaces are examples of types. Fields, methods, properties, and events are examples of members. When C# programs are compiled, they are physically packaged into assemblies. Assemblies typically have the file extension `.exe` or `.dll`, depending on whether they implement *applications* or *libraries*.

The example

```
using System;
namespace Acme.Collections
{
    public class Stack
    {
        Entry top;

        public void Push(object data) {
            top = new Entry(top, data);
        }

        public object Pop() {
            if (top == null) throw new InvalidOperationException();
            object result = top.data;
            top = top.next;
            return result;
        }
    }

    class Entry
    {
        public Entry next;
        public object data;
    }
}
```

```

        public Entry(Entry next, object data) {
            this.next = next;
            this.data = data;
        }
    }
}

```

declares a class named `Stack` in a namespace called `Acme.Collections`. The fully qualified name of this class is `Acme.Collections.Stack`. The class contains several members: a field named `top`, two methods named `Push` and `Pop`, and a nested class named `Entry`. The `Entry` class further contains three members: a field named `next`, a field named `data`, and a constructor. Assuming that the source code of the example is stored in the file `acme.cs`, the command line

```
csc /t:library acme.cs
```

compiles the example as a library (code without a `Main` entry point) and produces an assembly named `acme.dll`.

Assemblies contain executable code in the form of *Intermediate Language* (IL) instructions, and symbolic information in the form of *metadata*. Before it is executed, the IL code in an assembly is automatically converted to processor-specific code by the Just-In-Time (JIT) compiler of .NET Common Language Runtime.

Because an assembly is a self-describing unit of functionality containing both code and metadata, there is no need for `#include` directives and header files in C#. The public types and members contained in a particular assembly are made available in a C# program simply by referencing that assembly when compiling the program. For example, this program uses the `Acme.Collections.Stack` class from the `acme.dll` assembly:

```

using System;
using Acme.Collections;
class Test
{
    static void Main() {
        Stack s = new Stack();
        s.Push(1);
        s.Push(10);
        s.Push(100);
        Console.WriteLine(s.Pop());
        Console.WriteLine(s.Pop());
        Console.WriteLine(s.Pop());
    }
}

```

If the program is stored in the file `test.cs`, when `test.cs` is compiled, the `acme.dll` assembly can be referenced using the compiler's `/r` option:

```
csc /r:acme.dll test.cs
```

This creates an executable assembly named `test.exe`, which, when run, produces the output:

```

100
10
1

```

C# permits the source text of a program to be stored in several source files. When a multi-file C# program is compiled, all of the source files are processed together, and the source files can freely reference each other—conceptually, it is as if all the source files were concatenated into one large file before being processed. Forward declarations are never needed in C# because, with very few exceptions, declaration order is insignificant. C# does not limit a source file to declaring only one public type nor does it require the name of the source file to match a type declared in the source file.

1.3 Types and variables

There are two kinds of types in C#: *value types* and *reference types*. Variables of value types directly contain their data whereas variables of reference types store references to their data, the latter being known as objects. With reference types, it is possible for two variables to reference the same object and thus possible for operations on one variable to affect the object referenced by the other variable. With value types, the variables each have their own copy of the data, and it is not possible for operations on one to affect the other (except in the case of `ref` and `out` parameter variables).

C#'s value types are further divided into *simple types*, *enum types*, *struct types*, and *nullable types*, and C#'s reference types are further divided into *class types*, *interface types*, *array types*, and *delegate types*.

The following table provides an overview of C#'s type system.

Category		Description
Value types	Simple types	Signed integral: <code>sbyte</code> , <code>short</code> , <code>int</code> , <code>long</code>
		Unsigned integral: <code>byte</code> , <code>ushort</code> , <code>uint</code> , <code>ulong</code>
		Unicode characters: <code>char</code>
		IEEE floating point: <code>float</code> , <code>double</code>
		High-precision decimal: <code>decimal</code>
		Boolean: <code>bool</code>
	Enum types	User-defined types of the form <code>enum E { ... }</code>
Reference types	Struct types	User-defined types of the form <code>struct S { ... }</code>
	Nullable types	Extensions of all other value types with a <code>null</code> value
	Class types	Ultimate base class of all other types: <code>object</code>
		Unicode strings: <code>string</code>
		User-defined types of the form <code>class C { ... }</code>
	Interface types	User-defined types of the form <code>interface I { ... }</code>
	Array types	Single- and multi-dimensional, for example, <code>int[]</code> and <code>int[,]</code>
	Delegate types	User-defined types of the form e.g. <code>delegate int D(...)</code>

The eight integral types provide support for 8-bit, 16-bit, 32-bit, and 64-bit values in signed or unsigned form.

The two floating point types, `float` and `double`, are represented using the 32-bit single-precision and 64-bit double-precision IEEE 754 formats.

The `decimal` type is a 128-bit data type suitable for financial and monetary calculations.

C#'s `bool` type is used to represent boolean values—values that are either `true` or `false`.

Character and string processing in C# uses Unicode encoding. The `char` type represents a UTF-16 code unit, and the `string` type represents a sequence of UTF-16 code units.

The following table summarizes C#'s numeric types.

Category	Bits	Type	Range/Precision
Signed integral	8	sbyte	-128...127
	16	short	-32,768...32,767
	32	int	-2,147,483,648...2,147,483,647
	64	long	-9,223,372,036,854,775,808...9,223,372,036,854,775,807
Unsigned integral	8	byte	0...255
	16	ushort	0...65,535
	32	uint	0...4,294,967,295
	64	ulong	0...18,446,744,073,709,551,615
Floating point	32	float	1.5×10^{-45} to 3.4×10^{38} , 7-digit precision
	64	double	5.0×10^{-324} to 1.7×10^{308} , 15-digit precision
Decimal	128	decimal	1.0×10^{-28} to 7.9×10^{28} , 28-digit precision

C# programs use **type declarations** to create new types. A type declaration specifies the name and the members of the new type. Five of C#'s categories of types are user-definable: class types, struct types, interface types, enum types, and delegate types.

A class type defines a data structure that contains data members (fields) and function members (methods, properties, and others). Class types support single inheritance and polymorphism, mechanisms whereby derived classes can extend and specialize base classes.

A struct type is similar to a class type in that it represents a structure with data members and function members. However, unlike classes, structs are value types and do not require heap allocation. Struct types do not support user-specified inheritance, and all struct types implicitly inherit from type `object`.

An interface type defines a contract as a named set of public function members. A class or struct that implements an interface must provide implementations of the interface's function members. An interface may inherit from multiple base interfaces, and a class or struct may implement multiple interfaces.

A delegate type represents references to methods with a particular parameter list and return type. Delegates make it possible to treat methods as entities that can be assigned to variables and passed as parameters. Delegates are similar to the concept of function pointers found in some other languages, but unlike function pointers, delegates are object-oriented and type-safe.

Class, struct, interface and delegate types all support generics, whereby they can be parameterized with other types.

An enum type is a distinct type with named constants. Every enum type has an underlying type, which must be one of the eight integral types. The set of values of an enum type is the same as the set of values of the underlying type.

C# supports single- and multi-dimensional arrays of any type. Unlike the types listed above, array types do not have to be declared before they can be used. Instead, array types are constructed by following a type name with square brackets. For example, `int[]` is a single-dimensional array of `int`, `int[,]` is a two-dimensional array of `int`, and `int[][]` is a single-dimensional array of single-dimensional arrays of `int`.

Nullable types also do not have to be declared before they can be used. For each non-nullable value type `T` there is a corresponding nullable type `T?`, which can hold an additional value `null`. For instance, `int?` is a type that can hold any 32 bit integer or the value `null`.

C#'s type system is unified such that a value of any type can be treated as an object. Every type in C# directly or indirectly derives from the `object` class type, and `object` is the ultimate base class of all types. Values of reference types are treated as objects simply by viewing the values as type `object`. Values of value types are treated as objects by performing **boxing** and **unboxing** operations. In the following example, an `int` value is converted to `object` and back again to `int`.

```
using System;
class Test
{
    static void Main() {
        int i = 123;
        object o = i;      // Boxing
        int j = (int)o;    // Unboxing
    }
}
```

When a value of a value type is converted to type `object`, an object instance, also called a "box," is allocated to hold the value, and the value is copied into that box. Conversely, when an `object` reference is cast to a value type, a check is made that the referenced object is a box of the correct value type, and, if the check succeeds, the value in the box is copied out.

C#'s unified type system effectively means that value types can become objects "on demand." Because of the unification, general-purpose libraries that use type `object` can be used with both reference types and value types.

There are several kinds of **variables** in C#, including fields, array elements, local variables, and parameters. Variables represent storage locations, and every variable has a type that determines what values can be stored in the variable, as shown by the following table.

Type of Variable	Possible Contents
Non-nullable value type	A value of that exact type
Nullable value type	A null value or a value of that exact type
<code>object</code>	A null reference, a reference to an object of any reference type, or a reference to a boxed value of any value type
Class type	A null reference, a reference to an instance of that class type, or a reference to an instance of a class derived from that class type
Interface type	A null reference, a reference to an instance of a class type that implements that interface type, or a reference to a boxed value of a value type that implements that interface type
Array type	A null reference, a reference to an instance of that array type, or a reference to an instance of a compatible array type
Delegate type	A null reference or a reference to an instance of that delegate type

1.4 Expressions

Expressions are constructed from **operands** and **operators**. The operators of an expression indicate which operations to apply to the operands. Examples of operators include `+`, `-`, `*`, `/`, and `new`. Examples of operands include literals, fields, local variables, and expressions.

Chapter 1 Introduction

When an expression contains multiple operators, the *precedence* of the operators controls the order in which the individual operators are evaluated. For example, the expression `x + y * z` is evaluated as `x + (y * z)` because the `*` operator has higher precedence than the `+` operator.

Most operators can be *overloaded*. Operator overloading permits user-defined operator implementations to be specified for operations where one or both of the operands are of a user-defined class or struct type.

The following table summarizes C#'s operators, listing the operator categories in order of precedence from highest to lowest. Operators in the same category have equal precedence.

Category	Expression	Description
Primary	<code>x.m</code>	Member access
	<code>x(...)</code>	Method and delegate invocation
	<code>x[...]</code>	Array and indexer access
	<code>x++</code>	Post-increment
	<code>x--</code>	Post-decrement
	<code>new T(...)</code>	Object and delegate creation
	<code>new T(...){...}</code>	Object creation with initializer
	<code>new {...}</code>	Anonymous object initializer
	<code>new T[...]</code>	Array creation
	<code>typeof(T)</code>	Obtain <code>System.Type</code> object for <code>T</code>
	<code>checked(x)</code>	Evaluate expression in checked context
	<code>unchecked(x)</code>	Evaluate expression in unchecked context
	<code>default(T)</code>	Obtain default value of type <code>T</code>
	<code>delegate {...}</code>	Anonymous function (anonymous method)
Unary	<code>+x</code>	Identity
	<code>-x</code>	Negation
	<code>!x</code>	Logical negation
	<code>~x</code>	Bitwise negation
	<code>++x</code>	Pre-increment
	<code>--x</code>	Pre-decrement
	<code>(T)x</code>	Explicitly convert <code>x</code> to type <code>T</code>
Multiplicative	<code>x * y</code>	Multiplication
	<code>x / y</code>	Division
	<code>x % y</code>	Remainder
Additive	<code>x + y</code>	Addition, string concatenation, delegate combination
	<code>x - y</code>	Subtraction, delegate removal

C# LANGUAGE SPECIFICATION

Shift	<code>x << y</code>	Shift left
	<code>x >> y</code>	Shift right
Relational and type testing	<code>x < y</code>	Less than
	<code>x > y</code>	Greater than
	<code>x <= y</code>	Less than or equal
	<code>x >= y</code>	Greater than or equal
	<code>x is T</code>	Return <code>true</code> if <code>x</code> is a <code>T</code> , <code>false</code> otherwise
	<code>x as T</code>	Return <code>x</code> typed as <code>T</code> , or <code>null</code> if <code>x</code> is not a <code>T</code>
Equality	<code>x == y</code>	Equal
	<code>x != y</code>	Not equal
Logical AND	<code>x & y</code>	Integer bitwise AND, boolean logical AND
Logical XOR	<code>x ^ y</code>	Integer bitwise XOR, boolean logical XOR
Logical OR	<code>x y</code>	Integer bitwise OR, boolean logical OR
Conditional AND	<code>x && y</code>	Evaluates <code>y</code> only if <code>x</code> is <code>true</code>
Conditional OR	<code>x y</code>	Evaluates <code>y</code> only if <code>x</code> is <code>false</code>
Null coalescing	<code>x ?? y</code>	Evaluates to <code>y</code> if <code>x</code> is <code>null</code> , to <code>x</code> otherwise
Conditional	<code>x ? y : z</code>	Evaluates <code>y</code> if <code>x</code> is <code>true</code> , <code>z</code> if <code>x</code> is <code>false</code>
Assignment or anonymous function	<code>x = y</code>	Assignment
	<code>x op= y</code>	Compound assignment; supported operators are <code>*=</code> <code>/=</code> <code>%=</code> <code>+=</code> <code>-=</code> <code><<=</code> <code>>>=</code> <code>&=</code> <code>^=</code> <code> =</code>
	<code>(T x) => y</code>	Anonymous function (lambda expression)

1.5 Statements

The actions of a program are expressed using *statements*. C# supports several different kinds of statements, a number of which are defined in terms of embedded statements.

A *block* permits multiple statements to be written in contexts where a single statement is allowed. A block consists of a list of statements written between the delimiters `{` and `}`.

Declaration statements are used to declare local variables and constants.

Expression statements are used to evaluate expressions. Expressions that can be used as statements include method invocations, object allocations using the `new` operator, assignments using `=` and the compound assignment operators, and increment and decrement operations using the `++` and `--` operators.

Selection statements are used to select one of a number of possible statements for execution based on the value of some expression. In this group are the `if` and `switch` statements.

Iteration statements are used to repeatedly execute an embedded statement. In this group are the `while`, `do`, `for`, and `foreach` statements.

Chapter 1 Introduction

Jump statements are used to transfer control. In this group are the `break`, `continue`, `goto`, `throw`, `return`, and `yield` statements.

The `try...catch` statement is used to catch exceptions that occur during execution of a block, and the `try...finally` statement is used to specify finalization code that is always executed, whether an exception occurred or not.

The `checked` and `unchecked` statements are used to control the overflow checking context for integral-type arithmetic operations and conversions.

The `lock` statement is used to obtain the mutual-exclusion lock for a given object, execute a statement, and then release the lock.

The `using` statement is used to obtain a resource, execute a statement, and then dispose of that resource.

The following table lists C#'s statements and provides an example for each one.

Statement	Example
Local variable declaration	<pre>static void Main() { int a; int b = 2, c = 3; a = 1; Console.WriteLine(a + b + c); }</pre>
Local constant declaration	<pre>static void Main() { const float pi = 3.1415927f; const int r = 25; Console.WriteLine(pi * r * r); }</pre>
Expression statement	<pre>static void Main() { int i; i = 123; // Expression statement Console.WriteLine(i); // Expression statement i++; // Expression statement Console.WriteLine(i); // Expression statement }</pre>
if statement	<pre>static void Main(string[] args) { if (args.Length == 0) { Console.WriteLine("No arguments"); } else { Console.WriteLine("One or more arguments"); } }</pre>

C# LANGUAGE SPECIFICATION

switch statement	<pre>static void Main(string[] args) { int n = args.Length; switch (n) { case 0: Console.WriteLine("No arguments"); break; case 1: Console.WriteLine("One argument"); break; default: Console.WriteLine("{0} arguments", n); break; } }</pre>
while statement	<pre>static void Main(string[] args) { int i = 0; while (i < args.Length) { Console.WriteLine(args[i]); i++; } }</pre>
do statement	<pre>static void Main() { string s; do { s = Console.ReadLine(); if (s != null) Console.WriteLine(s); } while (s != null); }</pre>
for statement	<pre>static void Main(string[] args) { for (int i = 0; i < args.Length; i++) { Console.WriteLine(args[i]); } }</pre>
foreach statement	<pre>static void Main(string[] args) { foreach (string s in args) { Console.WriteLine(s); } }</pre>
break statement	<pre>static void Main() { while (true) { string s = Console.ReadLine(); if (s == null) break; Console.WriteLine(s); } }</pre>
continue statement	<pre>static void Main(string[] args) { for (int i = 0; i < args.Length; i++) { if (args[i].StartsWith("/")) continue; Console.WriteLine(args[i]); } }</pre>

goto statement	<pre>static void Main(string[] args) { int i = 0; goto check; loop: Console.WriteLine(args[i++]); check: if (i < args.Length) goto loop; }</pre>
return statement	<pre>static int Add(int a, int b) { return a + b; } static void Main() { Console.WriteLine(Add(1, 2)); return; }</pre>
yield statement	<pre>static IEnumerable<int> Range(int from, int to) { for (int i = from; i < to; i++) { yield return i; } yield break; } static void Main() { foreach (int x in Range(-10,10)) { Console.WriteLine(x); } }</pre>
throw and try statements	<pre>static double Divide(double x, double y) { if (y == 0) throw new DivideByZeroException(); return x / y; } static void Main(string[] args) { try { if (args.Length != 2) { throw new Exception("Two numbers required"); } double x = double.Parse(args[0]); double y = double.Parse(args[1]); Console.WriteLine(Divide(x, y)); } catch (Exception e) { Console.WriteLine(e.Message); } finally { Console.WriteLine("Good bye!"); } }</pre>
checked and unchecked statements	<pre>static void Main() { int i = int.MaxValue; checked { Console.WriteLine(i + 1); // Exception } unchecked { Console.WriteLine(i + 1); // Overflow } }</pre>

lock statement	<pre>class Account { decimal balance; public void withdraw(decimal amount) { lock (this) { if (amount > balance) { throw new Exception("Insufficient funds"); } balance -= amount; } } }</pre>
using statement	<pre>static void Main() { using (TextWriter w = File.CreateText("test.txt")) { w.WriteLine("Line one"); w.WriteLine("Line two"); w.WriteLine("Line three"); } }</pre>

1.6 Classes and objects

Classes are the most fundamental of C#'s types. A class is a data structure that combines state (fields) and actions (methods and other function members) in a single unit. A class provides a definition for dynamically created *instances* of the class, also known as *objects*. Classes support *inheritance* and *polymorphism*, mechanisms whereby *derived classes* can extend and specialize *base classes*.

New classes are created using class declarations. A class declaration starts with a header that specifies the attributes and modifiers of the class, the name of the class, the base class (if given), and the interfaces implemented by the class. The header is followed by the class body, which consists of a list of member declarations written between the delimiters { and }.

The following is a declaration of a simple class named Point:

```
public class Point
{
    public int x, y;
    public Point(int x, int y) {
        this.x = x;
        this.y = y;
    }
}
```

Instances of classes are created using the **new** operator, which allocates memory for a new instance, invokes a constructor to initialize the instance, and returns a reference to the instance. The following statements create two **Point** objects and store references to those objects in two variables:

```
Point p1 = new Point(0, 0);
Point p2 = new Point(10, 20);
```

The memory occupied by an object is automatically reclaimed when the object is no longer in use. It is neither necessary nor possible to explicitly deallocate objects in C#.

1.6.1 Members

The members of a class are either *static members* or *instance members*. Static members belong to classes, and instance members belong to objects (instances of classes).

The following table provides an overview of the kinds of members a class can contain.

Member	Description
Constants	Constant values associated with the class
Fields	Variables of the class
Methods	Computations and actions that can be performed by the class
Properties	Actions associated with reading and writing named properties of the class
Indexers	Actions associated with indexing instances of the class like an array
Events	Notifications that can be generated by the class
Operators	Conversions and expression operators supported by the class
Constructors	Actions required to initialize instances of the class or the class itself
Destructors	Actions to perform before instances of the class are permanently discarded
Types	Nested types declared by the class

1.6.2 Accessibility

Each member of a class has an associated accessibility, which controls the regions of program text that are able to access the member. There are five possible forms of accessibility. These are summarized in the following table.

Accessibility	Meaning
public	Access not limited
protected	Access limited to this class or classes derived from this class
internal	Access limited to this program
protected internal	Access limited to this program or classes derived from this class
private	Access limited to this class

1.6.3 Type parameters

A class definition may specify a set of type parameters by following the class name with angle brackets enclosing a list of type parameter names. The type parameters can be used in the body of the class declarations to define the members of the class. In the following example, the type parameters of `Pair` are `TFirst` and `TSecond`:

```
public class Pair<TFirst,TSecond>
{
    public TFirst First;
    public TSecond Second;
}
```

A class type that is declared to take type parameters is called a generic class type. Struct, interface and delegate types can also be generic.

When the generic class is used, type arguments must be provided for each of the type parameters:

```
Pair<int,string> pair = new Pair<int,string> { First = 1, Second = "two" };
int i = pair.First;    // TFirst is int
string s = pair.Second; // TSecond is string
```

A generic type with type arguments provided, like `Pair<int, string>` above, is called a constructed type.

1.6.4 Base classes

A class declaration may specify a base class by following the class name and type parameters with a colon and the name of the base class. Omitting a base class specification is the same as deriving from type `object`. In the following example, the base class of `Point3D` is `Point`, and the base class of `Point` is `object`:

```
public class Point
{
    public int x, y;
    public Point(int x, int y) {
        this.x = x;
        this.y = y;
    }
}

public class Point3D: Point
{
    public int z;
    public Point3D(int x, int y, int z): base(x, y) {
        this.z = z;
    }
}
```

A class inherits the members of its base class. Inheritance means that a class implicitly contains all members of its base class, except for the constructors of the base class. A derived class can add new members to those it inherits, but it cannot remove the definition of an inherited member. In the previous example, `Point3D` inherits the `x` and `y` fields from `Point`, and every `Point3D` instance contains three fields, `x`, `y`, and `z`.

An implicit conversion exists from a class type to any of its base class types. Therefore, a variable of a class type can reference an instance of that class or an instance of any derived class. For example, given the previous class declarations, a variable of type `Point` can reference either a `Point` or a `Point3D`:

```
Point a = new Point(10, 20);
Point b = new Point3D(10, 20, 30);
```

1.6.5 Fields

A field is a variable that is associated with a class or with an instance of a class.

A field declared with the `static` modifier defines a **static field**. A static field identifies exactly one storage location. No matter how many instances of a class are created, there is only ever one copy of a static field.

A field declared without the `static` modifier defines an **instance field**. Every instance of a class contains a separate copy of all the instance fields of that class.

In the following example, each instance of the `Color` class has a separate copy of the `r`, `g`, and `b` instance fields, but there is only one copy of the `Black`, `White`, `Red`, `Green`, and `Blue` static fields:

```
public class Color
{
    public static readonly Color Black = new Color(0, 0, 0);
    public static readonly Color White = new Color(255, 255, 255);
    public static readonly Color Red = new Color(255, 0, 0);
    public static readonly Color Green = new Color(0, 255, 0);
    public static readonly Color Blue = new Color(0, 0, 255);
    private byte r, g, b;
```

```

public Color(byte r, byte g, byte b) {
    this.r = r;
    this.g = g;
    this.b = b;
}
}

```

As shown in the previous example, **read-only fields** may be declared with a `readonly` modifier. Assignment to a `readonly` field can only occur as part of the field's declaration or in a constructor in the same class.

1.6.6 Methods

A **method** is a member that implements a computation or action that can be performed by an object or class. **Static methods** are accessed through the class. **Instance methods** are accessed through instances of the class.

Methods have a (possibly empty) list of **parameters**, which represent values or variable references passed to the method, and a **return type**, which specifies the type of the value computed and returned by the method. A method's return type is `void` if it does not return a value.

Like types, methods may also have a set of type parameters, for which type arguments must be specified when the method is called. Unlike types, the type arguments can often be inferred from the arguments of a method call and need not be explicitly given.

The **signature** of a method must be unique in the class in which the method is declared. The signature of a method consists of the name of the method, the number of type parameters and the number, modifiers, and types of its parameters. The signature of a method does not include the return type.

1.6.6.1 Parameters

Parameters are used to pass values or variable references to methods. The parameters of a method get their actual values from the **arguments** that are specified when the method is invoked. There are four kinds of parameters: value parameters, reference parameters, output parameters, and parameter arrays.

A **value parameter** is used for input parameter passing. A value parameter corresponds to a local variable that gets its initial value from the argument that was passed for the parameter. Modifications to a value parameter do not affect the argument that was passed for the parameter.

A **reference parameter** is used for both input and output parameter passing. The argument passed for a reference parameter must be a variable, and during execution of the method, the reference parameter represents the same storage location as the argument variable. A reference parameter is declared with the `ref` modifier. The following example shows the use of `ref` parameters.

```

using System;
class Test
{
    static void Swap(ref int x, ref int y) {
        int temp = x;
        x = y;
        y = temp;
    }

    static void Main() {
        int i = 1, j = 2;
        Swap(ref i, ref j);
        Console.WriteLine("{0} {1}", i, j);    // Outputs "2 1"
    }
}

```

An **output parameter** is used for output parameter passing. An output parameter is similar to a reference parameter except that the initial value of the caller-provided argument is unimportant. An output parameter is declared with the `out` modifier. The following example shows the use of `out` parameters.

```

using System;
class Test
{
    static void Divide(int x, int y, out int result, out int remainder) {
        result = x / y;
        remainder = x % y;
    }

    static void Main() {
        int res, rem;
        Divide(10, 3, out res, out rem);
        Console.WriteLine("{0} {1}", res, rem);    // Outputs "3 1"
    }
}

```

A **parameter array** permits a variable number of arguments to be passed to a method. A parameter array is declared with the `params` modifier. Only the last parameter of a method can be a parameter array, and the type of a parameter array must be a single-dimensional array type. The `write` and `writeLine` methods of the `System.Console` class are good examples of parameter array usage. They are declared as follows.

```

public class Console
{
    public static void Write(string fmt, params object[] args) {...}
    public static void WriteLine(string fmt, params object[] args) {...}
    ...
}

```

Within a method that uses a parameter array, the parameter array behaves exactly like a regular parameter of an array type. However, in an invocation of a method with a parameter array, it is possible to pass either a single argument of the parameter array type or any number of arguments of the element type of the parameter array. In the latter case, an array instance is automatically created and initialized with the given arguments. This example

```
Console.WriteLine("x={0} y={1} z={2}", x, y, z);
```

is equivalent to writing the following.

```

string s = "x={0} y={1} z={2}";
object[] args = new object[3];
args[0] = x;
args[1] = y;
args[2] = z;
Console.WriteLine(s, args);

```

1.6.6.2 Method body and local variables

A method's body specifies the statements to execute when the method is invoked.

A method body can declare variables that are specific to the invocation of the method. Such variables are called **local variables**. A local variable declaration specifies a type name, a variable name, and possibly an initial value. The following example declares a local variable `i` with an initial value of zero and a local variable `j` with no initial value.

```
using System;
```



```

class Squares
{
    static void Main() {
        int i = 0;
        int j;
        while (i < 10) {
            j = i * i;
            Console.WriteLine("{0} x {0} = {1}", i, j);
            i = i + 1;
        }
    }
}

```

C# requires a local variable to be *definitely assigned* before its value can be obtained. For example, if the declaration of the previous `i` did not include an initial value, the compiler would report an error for the subsequent usages of `i` because `i` would not be definitely assigned at those points in the program.

A method can use `return` statements to return control to its caller. In a method returning `void`, `return` statements cannot specify an expression. In a method returning non-`void`, `return` statements must include an expression that computes the return value.

1.6.6.3 Static and instance methods

A method declared with a `static` modifier is a *static method*. A static method does not operate on a specific instance and can only directly access static members.

A method declared without a `static` modifier is an *instance method*. An instance method operates on a specific instance and can access both static and instance members. The instance on which an instance method was invoked can be explicitly accessed as `this`. It is an error to refer to `this` in a static method.

The following `Entity` class has both static and instance members.

```

class Entity
{
    static int nextSerialNo;
    int serialNo;
    public Entity() {
        serialNo = nextSerialNo++;
    }
    public int GetSerialNo() {
        return serialNo;
    }
    public static int GetNextSerialNo() {
        return nextSerialNo;
    }
    public static void SetNextSerialNo(int value) {
        nextSerialNo = value;
    }
}

```

Each `Entity` instance contains a serial number (and presumably some other information that is not shown here). The `Entity` constructor (which is like an instance method) initializes the new instance with the next available serial number. Because the constructor is an instance member, it is permitted to access both the `serialNo` instance field and the `nextSerialNo` static field.

The `GetNextSerialNo` and `SetNextSerialNo` static methods can access the `nextSerialNo` static field, but it would be an error for them to directly access the `serialNo` instance field.

The following example shows the use of the `Entity` class.

```

using System;
class Test
{
    static void Main() {
        Entity.SetNextSerialNo(1000);
        Entity e1 = new Entity();
        Entity e2 = new Entity();

        Console.WriteLine(e1.GetSerialNo());           // Outputs "1000"
        Console.WriteLine(e2.GetSerialNo());           // Outputs "1001"
        Console.WriteLine(Entity.GetNextSerialNo());   // Outputs "1002"
    }
}

```

Note that the `SetNextSerialNo` and `GetNextSerialNo` static methods are invoked on the class whereas the `GetSerialNo` instance method is invoked on instances of the class.

1.6.6.4 Virtual, override, and abstract methods

When an instance method declaration includes a `virtual` modifier, the method is said to be a *virtual method*. When no `virtual` modifier is present, the method is said to be a *non-virtual method*.

When a virtual method is invoked, the *runtime type* of the instance for which that invocation takes place determines the actual method implementation to invoke. In a nonvirtual method invocation, the *compile-time type* of the instance is the determining factor.

A virtual method can be *overridden* in a derived class. When an instance method declaration includes an `override` modifier, the method overrides an inherited virtual method with the same signature. Whereas a virtual method declaration *introduces* a new method, an `override` method declaration *specializes* an existing inherited virtual method by providing a new implementation of that method.

An *abstract* method is a virtual method with no implementation. An abstract method is declared with the `abstract` modifier and is permitted only in a class that is also declared `abstract`. An abstract method must be overridden in every non-abstract derived class.

The following example declares an abstract class, `Expression`, which represents an expression tree node, and three derived classes, `Constant`, `VariableReference`, and `Operation`, which implement expression tree nodes for constants, variable references, and arithmetic operations. (This is similar to, but not to be confused with the expression tree types introduced in section §4.6).

```

using System;
using System.Collections;
public abstract class Expression
{
    public abstract double Evaluate(Hashtable vars);
}
public class Constant: Expression
{
    double value;
    public Constant(double value) {
        this.value = value;
    }
    public override double Evaluate(Hashtable vars) {
        return value;
    }
}

```

```

public class VariableReference: Expression
{
    string name;
    public VariableReference(string name) {
        this.name = name;
    }
    public override double Evaluate(Hashtable vars) {
        object value = vars[name];
        if (value == null) {
            throw new Exception("Unknown variable: " + name);
        }
        return Convert.ToDouble(value);
    }
}

public class Operation: Expression
{
    Expression left;
    char op;
    Expression right;
    public Operation(Expression left, char op, Expression right) {
        this.left = left;
        this.op = op;
        this.right = right;
    }
    public override double Evaluate(Hashtable vars) {
        double x = left.Evaluate(vars);
        double y = right.Evaluate(vars);
        switch (op) {
            case '+': return x + y;
            case '-': return x - y;
            case '*': return x * y;
            case '/': return x / y;
        }
        throw new Exception("Unknown operator");
    }
}

```

The previous four classes can be used to model arithmetic expressions. For example, using instances of these classes, the expression $x + 3$ can be represented as follows.

```

Expression e = new Operation(
    new VariableReference("x"),
    '+',
    new Constant(3));

```

The `Evaluate` method of an `Expression` instance is invoked to evaluate the given expression and produce a `double` value. The method takes as an argument a `Hashtable` that contains variable names (as keys of the entries) and values (as values of the entries). The `Evaluate` method is a virtual abstract method, meaning that non-abstract derived classes must override it to provide an actual implementation.

A `Constant`'s implementation of `Evaluate` simply returns the stored constant. A `VariableReference`'s implementation looks up the variable name in the hashtable and returns the resulting value. An `Operation`'s implementation first evaluates the left and right operands (by recursively invoking their `Evaluate` methods) and then performs the given arithmetic operation.

The following program uses the `Expression` classes to evaluate the expression $x * (y + 2)$ for different values of x and y .

```

using System;
using System.Collections;

```

```

class Test
{
    static void Main() {
        Expression e = new Operation(
            new VariableReference("x"),
            '*',
            new Operation(
                new VariableReference("y"),
                '+',
                new Constant(2)
            )
        );
        Hashtable vars = new Hashtable();
        vars["x"] = 3;
        vars["y"] = 5;
        Console.WriteLine(e.Evaluate(vars));    // Outputs "21"
        vars["x"] = 1.5;
        vars["y"] = 9;
        Console.WriteLine(e.Evaluate(vars));    // Outputs "16.5"
    }
}

```

1.6.6.5 Method overloading

Method **overloading** permits multiple methods in the same class to have the same name as long as they have unique signatures. When compiling an invocation of an overloaded method, the compiler uses **overload resolution** to determine the specific method to invoke. Overload resolution finds the one method that best matches the arguments or reports an error if no single best match can be found. The following example shows overload resolution in effect. The comment for each invocation in the `Main` method shows which method is actually invoked.

```

class Test
{
    static void F() {
        Console.WriteLine("F()");
    }
    static void F(object x) {
        Console.WriteLine("F(object)");
    }
    static void F(int x) {
        Console.WriteLine("F(int)");
    }
    static void F(double x) {
        Console.WriteLine("F(double)");
    }
    static void F<T>(T x) {
        Console.WriteLine("F<T>(T)");
    }
    static void F(double x, double y) {
        Console.WriteLine("F(double, double)");
    }
}

```

```

static void Main() {
    F();           // Invokes F()
    F(1);          // Invokes F(int)
    F(1.0);        // Invokes F(double)
    F("abc");      // Invokes F(object)
    F((double)1);  // Invokes F(double)
    F((object)1);  // Invokes F(object)
    F<int>(1);      // Invokes F<T>(T)
    F(1, 1);       // Invokes F(double, double) }
}

```

As shown by the example, a particular method can always be selected by explicitly casting the arguments to the exact parameter types and/or explicitly supplying type arguments.

1.6.7 Other function members

Members that contain executable code are collectively known as the *function members* of a class. The preceding section describes methods, which are the primary kind of function members. This section describes the other kinds of function members supported by C#: constructors, properties, indexers, events, operators, and destructors.

The following table shows a generic class called `List<T>`, which implements a growable list of objects. The class contains several examples of the most common kinds of function members.

<pre> public class List<T> { </pre>	
<pre> const int defaultCapacity = 4; </pre>	Constant
<pre> T[] items; int count; </pre>	Fields
<pre> public List(): this(defaultCapacity) {} public List(int capacity) { items = new T[capacity]; } </pre>	Constructors
<pre> public int Count { get { return count; } } public int Capacity { get { return items.Length; } set { if (value < count) value = count; if (value != items.Length) { T[] newItems = new T[value]; Array.Copy(items, 0, newItems, 0, count); items = newItems; } } } </pre>	Properties

<pre> public T this[int index] { get { return items[index]; } set { items[index] = value; OnChanged(); } } </pre>	Indexer
<pre> public void Add(T item) { if (count == Capacity) Capacity = count * 2; items[count] = item; count++; OnChanged(); } protected virtual void OnChanged() { if (Changed != null) Changed(this, EventArgs.Empty); } public override bool Equals(object other) { return Equals(this, other as List<T>); } static bool Equals(List<T> a, List<T> b) { if (a == null) return b == null; if (b == null a.count != b.count) return false; for (int i = 0; i < a.count; i++) { if (!object.Equals(a.items[i], b.items[i])) { return false; } } return true; } </pre>	Methods
<pre> public event EventHandler Changed; </pre>	Event
<pre> public static bool operator ==(List<T> a, List<T> b) { return Equals(a, b); } public static bool operator !=(List<T> a, List<T> b) { return !Equals(a, b); } </pre>	Operators
<pre> } </pre>	

1.6.7.1 Constructors

C# supports both instance and static constructors. An *instance constructor* is a member that implements the actions required to initialize an instance of a class. A *static constructor* is a member that implements the actions required to initialize a class itself when it is first loaded.

A constructor is declared like a method with no return type and the same name as the containing class. If a constructor declaration includes a `static` modifier, it declares a static constructor. Otherwise, it declares an instance constructor.

Instance constructors can be overloaded. For example, the `List<T>` class declares two instance constructors, one with no parameters and one that takes an `int` parameter. Instance constructors are invoked using the `new` operator. The following statements allocate two `List<String>` instances using each of the constructors of the `List` class.

```
List<string> list1 = new List<string>();
List<string> list2 = new List<string>(10);
```

Unlike other members, instance constructors are not inherited, and a class has no instance constructors other than those actually declared in the class. If no instance constructor is supplied for a class, then an empty one with no parameters is automatically provided.

1.6.7.2 Properties

Properties are a natural extension of fields. Both are named members with associated types, and the syntax for accessing fields and properties is the same. However, unlike fields, properties do not denote storage locations. Instead, properties have **accessors** that specify the statements to be executed when their values are read or written.

A property is declared like a field, except that the declaration ends with a **get** accessor and/or a **set** accessor written between the delimiters { and } instead of ending in a semicolon. A property that has both a **get** accessor and a **set** accessor is a **read-write property**, a property that has only a **get** accessor is a **read-only property**, and a property that has only a **set** accessor is a **write-only property**.

A **get** accessor corresponds to a parameterless method with a return value of the property type. Except as the target of an assignment, when a property is referenced in an expression, the **get** accessor of the property is invoked to compute the value of the property.

A **set** accessor corresponds to a method with a single parameter named **value** and no return type. When a property is referenced as the target of an assignment or as the operand of ++ or --, the **set** accessor is invoked with an argument that provides the new value.

The `List<T>` class declares two properties, `Count` and `Capacity`, which are read-only and read-write, respectively. The following is an example of use of these properties.

```
List<string> names = new List<string>();
names.Capacity = 100;           // Invokes set accessor
int i = names.Count;           // Invokes get accessor
int j = names.Capacity;         // Invokes get accessor
```

Similar to fields and methods, C# supports both instance properties and static properties. Static properties are declared with the `static` modifier, and instance properties are declared without it.

The accessor(s) of a property can be `virtual`. When a property declaration includes a `virtual`, `abstract`, or `override` modifier, it applies to the accessor(s) of the property.

1.6.7.3 Indexers

An **indexer** is a member that enables objects to be indexed in the same way as an array. An indexer is declared like a property except that the name of the member is `this` followed by a parameter list written between the delimiters [and]. The parameters are available in the accessor(s) of the indexer. Similar to properties, indexers can be read-write, read-only, and write-only, and the accessor(s) of an indexer can be `virtual`.

The `List` class declares a single read-write indexer that takes an `int` parameter. The indexer makes it possible to index `List` instances with `int` values. For example

```
List<string> names = new List<string>();
names.Add("Liz");
names.Add("Martha");
names.Add("Beth");
for (int i = 0; i < names.Count; i++) {
    string s = names[i];
    names[i] = s.ToUpper();
}
```

Indexers can be overloaded, meaning that a class can declare multiple indexers as long as the number or types of their parameters differ.

1.6.7.4 Events

An **event** is a member that enables a class or object to provide notifications. An event is declared like a field except that the declaration includes an **event** keyword and the type must be a delegate type.

Within a class that declares an event member, the event behaves just like a field of a delegate type (provided the event is not abstract and does not declare accessors). The field stores a reference to a delegate that represents the event handlers that have been added to the event. If no event handles are present, the field is `null`.

The `List<T>` class declares a single event member called `Changed`, which indicates that a new item has been added to the list. The `Changed` event is raised by the `OnChanged` virtual method, which first checks whether the event is `null` (meaning that no handlers are present). The notion of raising an event is precisely equivalent to invoking the delegate represented by the event—thus, there are no special language constructs for raising events.

Clients react to events through **event handlers**. Event handlers are attached using the += operator and removed using the -= operator. The following example attaches an event handler to the `Changed` event of a `List<string>`.

```
using System;
class Test
{
    static int changeCount;

    static void ListChanged(object sender, EventArgs e) {
        changeCount++;
    }

    static void Main() {
        List<string> names = new List<string>();
        names.Changed += new EventHandler(ListChanged);
        names.Add("Liz");
        names.Add("Martha");
        names.Add("Beth");
        Console.WriteLine(changeCount);    // Outputs "3"
    }
}
```

For advanced scenarios where control of the underlying storage of an event is desired, an event declaration can explicitly provide `add` and `remove` accessors, which are somewhat similar to the `set` accessor of a property.

1.6.7.5 Operators

An **operator** is a member that defines the meaning of applying a particular expression operator to instances of a class. Three kinds of operators can be defined: unary operators, binary operators, and conversion operators. All operators must be declared as `public` and `static`.

The `List<T>` class declares two operators, `operator ==` and `operator !=`, and thus gives new meaning to expressions that apply those operators to `List` instances. Specifically, the operators define equality of two `List<T>` instances as comparing each of the contained objects using their `Equals` methods. The following example uses the `==` operator to compare two `List<int>` instances.

```
using System;
```

```

class Test
{
    static void Main() {
        List<int> a = new List<int>();
        a.Add(1);
        a.Add(2);
        List<int> b = new List<int>();
        b.Add(1);
        b.Add(2);
        Console.WriteLine(a == b);    // Outputs "True"
        a.Add(3);
        Console.WriteLine(a == b);    // Outputs "False"
    }
}

```

The first `Console.WriteLine` outputs `True` because the two lists contain the same number of objects with the same values in the same order. Had `List<T>` not defined operator `==`, the first `Console.WriteLine` would have output `False` because `a` and `b` reference different `List<int>` instances.

1.6.7.6 Destructors

A **destructor** is a member that implements the actions required to destruct an instance of a class. Destructors cannot have parameters, they cannot have accessibility modifiers, and they cannot be invoked explicitly. The destructor for an instance is invoked automatically during garbage collection.

The garbage collector is allowed wide latitude in deciding when to collect objects and run destructors. Specifically, the timing of destructor invocations is not deterministic, and destructors may be executed on any thread. For these and other reasons, classes should implement destructors only when no other solutions are feasible.

The `using` statement provides a better approach to object destruction.

1.7 Structs

Like classes, **structs** are data structures that can contain data members and function members, but unlike classes, structs are value types and do not require heap allocation. A variable of a struct type directly stores the data of the struct, whereas a variable of a class type stores a reference to a dynamically allocated object. Struct types do not support user-specified inheritance, and all struct types implicitly inherit from type `object`.

Structs are particularly useful for small data structures that have value semantics. Complex numbers, points in a coordinate system, or key-value pairs in a dictionary are all good examples of structs. The use of structs rather than classes for small data structures can make a large difference in the number of memory allocations an application performs. For example, the following program creates and initializes an array of 100 points. With `Point` implemented as a class, 101 separate objects are instantiated—one for the array and one each for the 100 elements.

```

class Point
{
    public int x, y;
    public Point(int x, int y) {
        this.x = x;
        this.y = y;
    }
}

```

```

class Test
{
    static void Main() {
        Point[] points = new Point[100];
        for (int i = 0; i < 100; i++) points[i] = new Point(i, i);
    }
}

```

An alternative is to make `Point` a struct.

```

struct Point
{
    public int x, y;
    public Point(int x, int y) {
        this.x = x;
        this.y = y;
    }
}

```

Now, only one object is instantiated—the one for the array—and the `Point` instances are stored in-line in the array.

Struct constructors are invoked with the `new` operator, but that does not imply that memory is being allocated. Instead of dynamically allocating an object and returning a reference to it, a struct constructor simply returns the struct value itself (typically in a temporary location on the stack), and this value is then copied as necessary.

With classes, it is possible for two variables to reference the same object and thus possible for operations on one variable to affect the object referenced by the other variable. With structs, the variables each have their own copy of the data, and it is not possible for operations on one to affect the other. For example, the output produced by the following code fragment depends on whether `Point` is a class or a struct.

```

Point a = new Point(10, 10);
Point b = a;
a.x = 20;
Console.WriteLine(b.x);

```

If `Point` is a class, the output is 20 because `a` and `b` reference the same object. If `Point` is a struct, the output is 10 because the assignment of `a` to `b` creates a copy of the value, and this copy is unaffected by the subsequent assignment to `a.x`.

The previous example highlights two of the limitations of structs. First, copying an entire struct is typically less efficient than copying an object reference, so assignment and value parameter passing can be more expensive with structs than with reference types. Second, except for `ref` and `out` parameters, it is not possible to create references to structs, which rules out their usage in a number of situations.

1.8 Arrays

An **array** is a data structure that contains a number of variables that are accessed through computed indices. The variables contained in an array, also called the **elements** of the array, are all of the same type, and this type is called the **element type** of the array.

Array types are reference types, and the declaration of an array variable simply sets aside space for a reference to an array instance. Actual array instances are created dynamically at runtime using the `new` operator. The `new` operation specifies the **length** of the new array instance, which is then fixed for the lifetime of the instance. The indices of the elements of an array range from 0 to `Length - 1`. The `new` operator automatically initializes the elements of an array to their default value, which, for example, is zero for all numeric types and `null` for all reference types.

The following example creates an array of `int` elements, initializes the array, and prints out the contents of the array.

```
using System;
class Test
{
    static void Main() {
        int[] a = new int[10];
        for (int i = 0; i < a.Length; i++) {
            a[i] = i * i;
        }
        for (int i = 0; i < a.Length; i++) {
            Console.WriteLine("a[{0}] = {1}", i, a[i]);
        }
    }
}
```

This example creates and operates on a *single-dimensional array*. C# also supports *multi-dimensional arrays*. The number of dimensions of an array type, also known as the *rank* of the array type, is one plus the number of commas written between the square brackets of the array type. The following example allocates a one-dimensional, a two-dimensional, and a three-dimensional array.

```
int[] a1 = new int[10];
int[,] a2 = new int[10, 5];
int[,,] a3 = new int[10, 5, 2];
```

The `a1` array contains 10 elements, the `a2` array contains 50 (10×5) elements, and the `a3` array contains 100 ($10 \times 5 \times 2$) elements.

The element type of an array can be any type, including an array type. An array with elements of an array type is sometimes called a *jagged array* because the lengths of the element arrays do not all have to be the same. The following example allocates an array of arrays of `int`:

```
int[][] a = new int[3][];
a[0] = new int[10];
a[1] = new int[5];
a[2] = new int[20];
```

The first line creates an array with three elements, each of type `int[]` and each with an initial value of `null`. The subsequent lines then initialize the three elements with references to individual array instances of varying lengths.

The `new` operator permits the initial values of the array elements to be specified using an *array initializer*, which is a list of expressions written between the delimiters `{` and `}`. The following example allocates and initializes an `int[]` with three elements.

```
int[] a = new int[] {1, 2, 3};
```

Note that the length of the array is inferred from the number of expressions between `{` and `}`. Local variable and field declarations can be shortened further such that the array type does not have to be restated.

```
int[] a = {1, 2, 3};
```

Both of the previous examples are equivalent to the following:

```
int[] t = new int[3];
t[0] = 1;
t[1] = 2;
t[2] = 3;
int[] a = t;
```

1.9 Interfaces

An *interface* defines a contract that can be implemented by classes and structs. An interface can contain methods, properties, events, and indexers. An interface does not provide implementations of the members it

defines—it merely specifies the members that must be supplied by classes or structs that implement the interface.

Interfaces may employ *multiple inheritance*. In the following example, the interface `IEnumerator` inherits from both `IEnumerator` and `IEnumerator`.

```
interface IEnumerator
{
    void Paint();
}

interface IEnumerator: IEnumerator
{
    void SetText(string text);
}

interface IEnumerator: IEnumerator
{
    void SetItems(string[] items);
}

interface IEnumerator: IEnumerator, IEnumerator {}
```

Classes and structs can implement multiple interfaces. In the following example, the class `EditBox` implements both `IEnumerator` and `IEnumerator`.

```
interface IEnumerator
{
    void Bind(Binder b);
}

public class EditBox: IEnumerator, IEnumerator
{
    public void Paint() {...}
    public void Bind(Binder b) {...}
}
```

When a class or struct implements a particular interface, instances of that class or struct can be implicitly converted to that interface type. For example

```
EditBox editBox = new EditBox();
IEnumerator control = editBox;
IEnumerator dataBound = editBox;
```

In cases where an instance is not statically known to implement a particular interface, dynamic type casts can be used. For example, the following statements use dynamic type casts to obtain an object's `IEnumerator` and `IEnumerator` interface implementations. Because the actual type of the object is `EditBox`, the casts succeed.

```
object obj = new EditBox();
IEnumerator control = (IEnumerator)obj;
IEnumerator dataBound = (IEnumerator)obj;
```

In the previous `EditBox` class, the `Paint` method from the `IEnumerator` interface and the `Bind` method from the `IEnumerator` interface are implemented using `public` members. C# also supports *explicit interface member implementations*, using which the class or struct can avoid making the members `public`. An explicit interface member implementation is written using the fully qualified interface member name. For example, the `EditBox` class could implement the `IEnumerator.Paint` and `IEnumerator.Bind` methods using explicit interface member implementations as follows.

```
public class EditBox: IEnumerator, IEnumerator
{
    void IEnumerator.Paint() {...}
    void IEnumerator.Bind(Binder b) {...}
}
```

Explicit interface members can only be accessed via the interface type. For example, the implementation of `IControl.Paint` provided by the previous `EditBox` class can only be invoked by first converting the `EditBox` reference to the `IControl` interface type.

```
EditBox editBox = new EditBox();
editBox.Paint();           // Error, no such method
IControl control = editBox;
control.Paint();           // ok
```

1.10 Enums

An *enum type* is a distinct value type with a set of named constants. The following example declares and uses an enum type named `Color` with three constant values, `Red`, `Green`, and `Blue`.

```
using System;
enum Color
{
    Red,
    Green,
    Blue
}
class Test
{
    static void PrintColor(Color color) {
        switch (color) {
            case Color.Red:
                Console.WriteLine("Red");
                break;
            case Color.Green:
                Console.WriteLine("Green");
                break;
            case Color.Blue:
                Console.WriteLine("Blue");
                break;
            default:
                Console.WriteLine("Unknown color");
                break;
        }
    }
    static void Main() {
        Color c = Color.Red;
        PrintColor(c);
        PrintColor(Color.Blue);
    }
}
```

Each enum type has a corresponding integral type called the *underlying type* of the enum type. An enum type that does not explicitly declare an underlying type has an underlying type of `int`. An enum type's storage format and range of possible values are determined by its underlying type. The set of values that an enum type can take on is not limited by its enum members. In particular, any value of the underlying type of an enum can be cast to the enum type and is a distinct valid value of that enum type.

The following example declares an enum type named `Alignment` with an underlying type of `sbyte`.

```
enum Alignment: sbyte
{
    Left = -1,
    Center = 0,
    Right = 1
}
```

As shown by the previous example, an enum member declaration can include a constant expression that specifies the value of the member. The constant value for each enum member must be in the range of the underlying type of the enum. When an enum member declaration does not explicitly specify a value, the member is given the value zero (if it is the first member in the enum type) or the value of the textually preceding enum member plus one.

Enum values can be converted to integral values and vice versa using type casts. For example

```
int i = (int)Color.Blue;           // int i = 2;
Color c = (Color)2;                // Color c = Color.Blue;
```

The default value of any enum type is the integral value zero converted to the enum type. In cases where variables are automatically initialized to a default value, this is the value given to variables of enum types. In order for the default value of an enum type to be easily available, the literal `0` implicitly converts to any enum type. Thus, the following is permitted.

```
Color c = 0;
```

1.11 Delegates

A *delegate type* represents references to methods with a particular parameter list and return type. Delegates make it possible to treat methods as entities that can be assigned to variables and passed as parameters. Delegates are similar to the concept of function pointers found in some other languages, but unlike function pointers, delegates are object-oriented and type-safe.

The following example declares and uses a delegate type named `Function`.

```
using System;
delegate double Function(double x);
class Multiplier
{
    double factor;
    public Multiplier(double factor) {
        this.factor = factor;
    }
    public double Multiply(double x) {
        return x * factor;
    }
}
class Test
{
    static double Square(double x) {
        return x * x;
    }
    static double[] Apply(double[] a, Function f) {
        double[] result = new double[a.Length];
        for (int i = 0; i < a.Length; i++) result[i] = f(a[i]);
        return result;
    }
    static void Main() {
        double[] a = {0.0, 0.5, 1.0};
        double[] squares = Apply(a, Square);
        double[] sines = Apply(a, Math.Sin);
        Multiplier m = new Multiplier(2.0);
        double[] doubles = Apply(a, m.Multiply);
    }
}
```

An instance of the `Function` delegate type can reference any method that takes a `double` argument and returns a `double` value. The `Apply` method applies a given `Function` to the elements of a `double[]`, returning a `double[]` with the results. In the `Main` method, `Apply` is used to apply three different functions to a `double[]`.

A delegate can reference either a static method (such as `Square` or `Math.Sin` in the previous example) or an instance method (such as `m.Multiply` in the previous example). A delegate that references an instance method also references a particular object, and when the instance method is invoked through the delegate, that object becomes `this` in the invocation.

Delegates can also be created using anonymous functions, which are “inline methods” that are created on the fly. Anonymous functions can see the local variables of the surrounding methods. Thus, the multiplier example above can be written more easily without using a `Multiplier` class:

```
double[] doubles = Apply(a, (double x) => x * 2.0);
```

An interesting and useful property of a delegate is that it does not know or care about the class of the method it references; all that matters is that the referenced method has the same parameters and return type as the delegate.

1.12 Attributes

Types, members, and other entities in a C# program support modifiers that control certain aspects of their behavior. For example, the accessibility of a method is controlled using the `public`, `protected`, `internal`, and `private` modifiers. C# generalizes this capability such that user-defined types of declarative information can be attached to program entities and retrieved at runtime. Programs specify this additional declarative information by defining and using *attributes*.

The following example declares a `HelpAttribute` attribute that can be placed on program entities to provide links to their associated documentation.

```
using System;
public class HelpAttribute: Attribute
{
    string url;
    string topic;

    public HelpAttribute(string url) {
        this.url = url;
    }

    public string Url {
        get { return url; }
    }

    public string Topic {
        get { return topic; }
        set { topic = value; }
    }
}
```

All attribute classes derive from the `System.Attribute` base class provided by the .NET Framework. Attributes can be applied by giving their name, along with any arguments, inside square brackets just before the associated declaration. If an attribute’s name ends in `Attribute`, that part of the name can be omitted when the attribute is referenced. For example, the `HelpAttribute` attribute can be used as follows.

```
[Help("http://msdn.microsoft.com/.../MyClass.htm")]
public class widget
{
    [Help("http://msdn.microsoft.com/.../MyClass.htm", Topic = "Display")]
    public void Display(string text) {}
}
```

This example attaches a `HelpAttribute` to the `Widget` class and another `HelpAttribute` to the `Display` method in the class. The public constructors of an attribute class control the information that must be provided when the attribute is attached to a program entity. Additional information can be provided by referencing public read-write properties of the attribute class (such as the reference to the `Topic` property previously).

The following example shows how attribute information for a given program entity can be retrieved at runtime using reflection.

```
using System;
using System.Reflection;
class Test
{
    static void ShowHelp(MemberInfo member) {
        HelpAttribute a = Attribute.GetCustomAttribute(member,
            typeof(HelpAttribute)) as HelpAttribute;
        if (a == null) {
            Console.WriteLine("No help for {0}", member);
        }
        else {
            Console.WriteLine("Help for {0}:", member);
            Console.WriteLine("  Url={0}, Topic={1}", a.Url, a.Topic);
        }
    }

    static void Main() {
        ShowHelp(typeof(Widget));
        ShowHelp(typeof(Widget).GetMethod("Display"));
    }
}
```

When a particular attribute is requested through reflection, the constructor for the attribute class is invoked with the information provided in the program source, and the resulting attribute instance is returned. If additional information was provided through properties, those properties are set to the given values before the attribute instance is returned.